

PDSA Notes

by

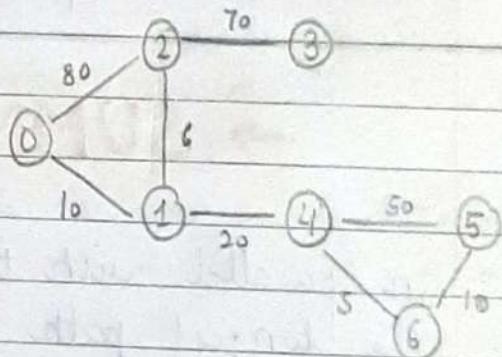
Gagreet Kaur

WEEK 5

Shortest Paths In Weighted Graphs

WEIGHTED GRAPHS

- Recall that BFS explores a graph level by level
- BFS computes shortest path , in terms of number of edges , to every reachable vertex
- May assign values to edges
 - cost , time , distance
 - Weighted graph
- $G = (V, E)$, $w : E \rightarrow \mathbb{R}$
- ADJACENCY MATRIX : Record weights along with edge information - weight is always zero if no edge
- ADJACENCY LIST : Record weights along with edge information



- In weighted graphs, each edge has a cost.
- Length of a path is the sum of weights

SHORTEST PATHS IN WEIGHTED GRAPHS

- BFS computes shortest path, in terms of number of edges, to every reachable vertex
- In a weighted graph, add up the weights along a path
- Weighted shortest path need not have minimum no. of edges
 - shortest path from 0 to 2 is via 1.

SHORTEST PATH PROBLEMS

SINGLE SOURCE SHORTEST PATHS

- Find shortest paths from a fixed vertex to every other vertex
- Transport finished product from factory (single source) to all retail outlets
- Courier company delivers items from distribution centre (single source) to addresses

ALL PAIRS SHORTEST PATHS

- Find shortest paths between every pair of vertices i & j
- Optimal airline, railway, road routes between cities

Negative Edge Weights

- should not have negative cycles
- w/o negative cycles, shortest paths still well defined

Chittenden

PAGE NO.

DATE:

NEGATIVE EDGE WEIGHTS

- Can negative edge weights be meaningful?
- Taxi driver trying to head home at the end of the day.
 - Roads with few customers, drive empty (the weight)
 - Roads with many customers, make profit (-ve weight)
 - Find a route toward home that minimizes the cost

NEGATIVE CYCLES

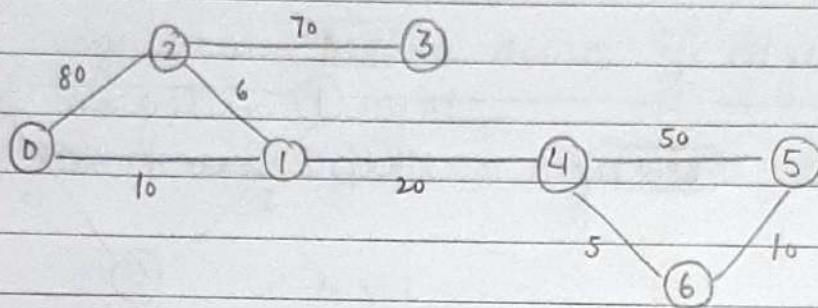
- A negative cycle is one whose weight is negative
 - sum of the weights of edges that make up the cycle
- By repeatedly traversing a negative cycle, total cost keeps decreasing
- If a graph has a negative cycle, shortest paths are not defined
- Without negative cycles, we can compute shortest paths even if some weights are negative

Date
January 25, 2022

Chitrang
PAGE NO.
DATE:

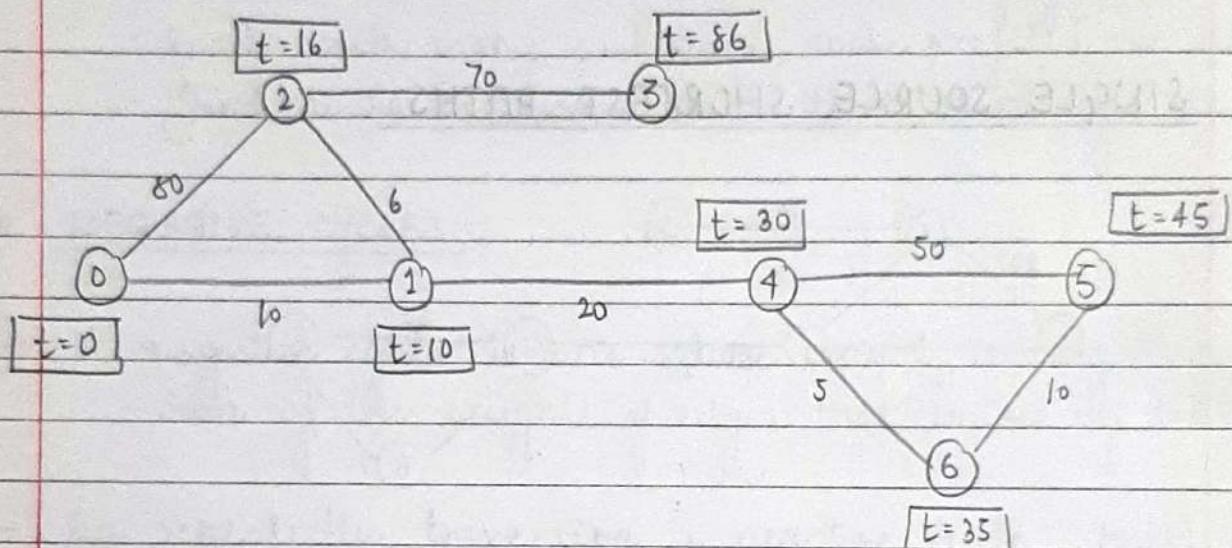
Single Source Shortest Paths

SINGLE SOURCE SHORTEST PATHS



- Weighted graph: $G = (V, E)$
 $w: E \rightarrow \mathbb{R}$
- Single Source Shortest Paths
 - find shortest paths from a fixed vertex to every other vertex
- Assume, for now, that edge weights are all non-negative
- Compute shortest paths from 0 to all the other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0

- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
-



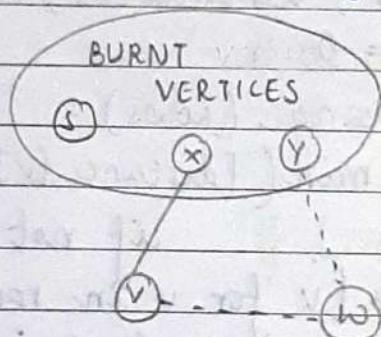
- Compute 'expected burn time' for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours
- Algorithm due to **Edsger W Dijkstra**

DIJKSTRA'S ALGORITHM : proof of correctness

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths

to all vertices already burnt.

- Next vertex to burn is v , via x
- cannot find a shorter path later from y to v via w
 - Burn time of $w \geq$ burn time of v
 - Edge from w to v has weight ≥ 0
- This argument breaks down if edge (w, v) can have negative weight
 - Can't use Dijkstra's Algo with negative edge wts.



IMPLEMENTATION

- Maintain 2 dictionaries with vertices as keys
 - visited, initially False for all v (burnt vertices)
 - distance, initially infinity for all v (expected burn time)
- Set distance [s] to 0
- Repeat until all reachable vertices are visited
 - Find unvisited vertex nextv with min. distance

- set `visited[nextv]` to True
- Recompute `distance[v]` for every neighbour v of `nextv`

CODE → : (1) Using Adjacency Matrix

```

def dijkstra (WMat, s):
    [rows, cols, x] = WMat.shape
    infinity = np.max (WMat) * rows + 1
    initialization
    (visited, distance) = ({}, {})
    for v in range (rows):
        (visited[v], distance[v]) = (False, infinity)
    distance[s] = 0
    for u in range (rows):
        nextd = min ([distance[v] for v in range (rows)
                     if not visited[v]])
        nextvlist = [v for v in range (rows)
                     if (not visited[v]) and
                        distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min (nextvlist)
        visited[nextv] = True
        for v in range (cols):
            if WMat[nextv, v, 0] == 1 and (not visited[v]):
                distance[v] = min (distance[v], distance[nextv]
                                   + WMat[nextv, v, 1])
return (distance)

```

(2) Using Adjacency List

```

def dijkstralist (wlist, s):
    infinity = 1 + len (wlist.keys ()) *
        max ([d for u in wlist.keys () for (v,d) in wlist[u]])
    (visited, distance) = ({}, {})
    for v in wlist.keys () :
        (visited[v], distance[v]) = (False, infinity)
    distance[s] = 0
    for u in wlist.keys () :
        nextd = min ([distance[v] for v in wlist.keys () if not visited[v]])
        nextvlist = [v for v in wlist.keys () if (not visited[v]) and distance[v] == nextd]
        if nextvlist == [] :
            break
        nextv = min (nextvlist)
        visited[nextv] = True
        for (v,d) in wlist[nextv] :
            if not visited[v] :
                distance[v] = min (distance[v], distance[nextv] + d)
    return (distance)

```

COMPLEXITY

- setting infinity takes $O(n^2)$ time
- Main loop runs n times
 - Each iteration visits one more vertex
 - $O(n)$ to find next vertex to visit
 - $O(n)$ to update $\text{distance}[v]$ for neighbours
- Overall $O(n^2)$
- If we use an adjacency list
 - Setting infinity & updating distances both $O(m)$, amortized
 - $O(n)$ bottleneck remains to find next vertex to visit
 - Better data structure?
- Dijkstra's algorithm computes single source shortest paths
- Use a greedy strategy to identify vertices to visit
 - Next vertex to visit is based on shortest distance computed so far
 - Need to prove that such a strategy is correct
 - Correctness requires edge weights to be non-negative
- Complexity is $O(n^2)$
 - Even with adj. list
 - Bottleneck is identifying unvisited vertex with min. distance

Single Source Shortest Paths with Negative Wts.

DIJKSTRA'S ALGORITHM

- Recall the burning pipeline analogy
- We keep track of the following.
 - The vertices that have been burnt
 - The expected burn time of vertices
- Initially.
 - No vertex is burnt
 - Expected burn time of source vertex is 0
 - Expected burn time of rest is ∞
- While there are vertices yet to burn
 - Pick unburnt vertex with minimum expected burn time, mark it as burnt
 - Update the expected burn time of its neighbours

INITIALIZATION (assume source vertex 0)

- $B(i) = \text{False}$, for $0 \leq i < n$
- $UB = \{k \mid B(k) = \text{False}\}$

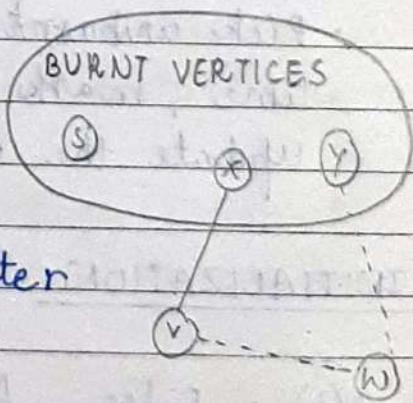
$$\cdot EBT(i) = \begin{cases} 0 & , \text{ if } i = 0 \\ \infty & ; \text{ otherwise} \end{cases}$$

UPDATE, if $UB \neq \emptyset$

- Let $j \in UB$ such that $EBT(j) \leq EBT(k)$ for all $k \in UB$
- Update $B(j) = \text{True}$, $UB = UB \setminus \{j\}$
- For each $(j, k) \in E$ such that $k \in UB$,
 $EBT(k) = \min(EBT(k), EBT(j) + w(j, k))$

CORRECTNESS REQUIRES NON-NEGATIVE EDGE WEIGHTS

- Each new shortest path we discover extends an earlier one
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is v , via x
- Cannot find a shorter path later from y to v via w
 - Burn time of $w \geq$ burn time of v
 - Edge from w to v has weight ≥ 0



- This argument breaks down if edge (w, v) can have negative weight

EXTENDING TO NEGATIVE EDGE WEIGHTS

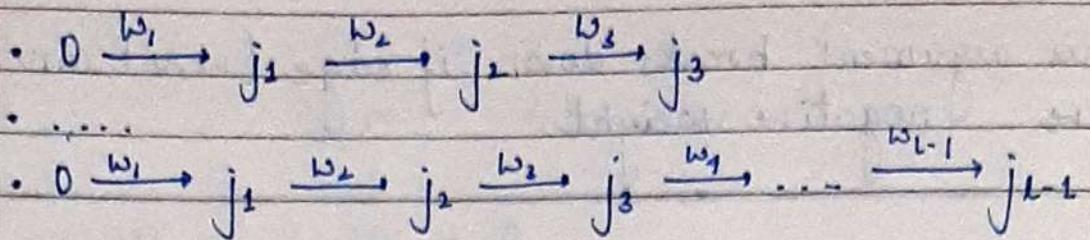
- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?
- Recall, negative edge weights are allowed, but no negative cycles
- Going around a cycle can only add to the length
- Shortest route to every vertex is a path, no loops
- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{l-1}} j_{l-1} \xrightarrow{w_l} k$$

Need not be minimum in terms of no. of edges

- Every prefix of this path must itself be a minimum weight path

$$\begin{aligned} & 0 \xrightarrow{w_1} j_1 \\ & 0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \end{aligned}$$



- Once we discover shortest path to j_{l-1} , next update will fix shortest path to k
- Repeatedly update shortest distance to each vertex based on shortest distance to its neighbours
 - Update can't push this distance below actual shortest distance
- After l updates, all shortest paths using $\leq l$ edges have stabilized
 - Minimum weight path to any node has at most $n-1$ edges
 - After $n-1$ updates, all shortest paths have stabilized

BELLMAN - FORD ALGORITHM

INITIALIZATION (source vertex 0)

$D(j)$: minimum distance known so far to vertex j

$$D(j) = \begin{cases} 0 & , \text{ if } j = 0 \\ \infty & , \text{ otherwise} \end{cases}$$

Repeat $n-1$ times

for each vertex $j \in \{0, 1, \dots, n-1\}$ for each edge $(j, k) \in E$,

$$D(k) = \min(D(k), D(j) + w(j, k))$$

works for directed & undirected graphs

WDE : (i) USING ADJACENCY MATRIX

def bellmanford (wMat, s) :

(rows, cols, s) = wMat.shape

infinity = np. max (wMat) * rows + 1

distance = {}

for v in range (rows) :

 distance [v] = infinity

 distance [s] = 0

 for i in range (rows) :

 for u in range (rows) :

 for v in range (cols) :

 if wMat [u, v, 0] == 1 :

 distance [v] = min (distance [v],

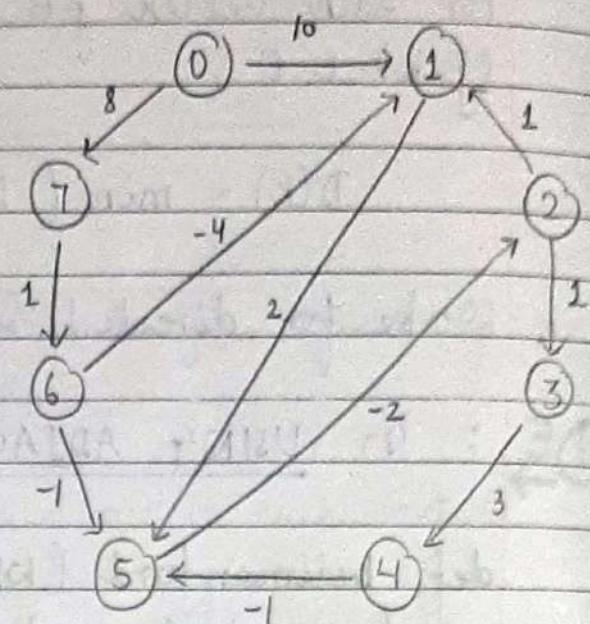
 distance [u] + wMat [u, v, 1])

return (distance)

Initialize $D(0) = 0$

For each $(j, k) \in E$, update $D(k) = \min(D(k), D(j) + w(j, k))$

v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



What if there was a negative cycle?

Distance would continue to decrease

Check if update v reduces $D(v)$ ^{any}

COMPLEXITY

- Initializing infinity takes $O(n^2)$ time
- The outer update loop runs $O(n)$ times
- In each iteration, we have to examine every edge in the graph
 - This take $O(n^2)$ for an adjacency matrix
- Overall, $O(n^3)$

→ If we shift to ADJACENCY LISTS

- Initializing infinity is $O(m)$

- Scanning all edges in each update iteration is $O(m)$

CODE : (2) USING ADJACENCY LIST

```
def bellmanfordlist (wList, s):
    infinity = 1 + len (wList.keys ()) *
        max ([d for u in wList.keys ()
              for (v, d) in wList [u]])
```

distance = {}

for v in wList.keys () :

distance [v] = infinity

distance [s] = 0

for i in range (n) :

for u in wList.keys () :

for (v, d) in wList [u] :

distance [v] = min (distance [v],

distance [u] + d)

return (distance)

→ Now, overall $O(mn)$

All-Pairs Shortest Paths

SHORTEST PATHS IN WEIGHTED GRAPHS

- Two types of shortest path problems of interest

SINGLE SOURCE SHORTEST PATHS

- Find shortest paths from a fixed vertex to every other vertex
- Transport finished product from factory (single source) to all retail outlets
- Courier company delivers items from distribution centre (single source) to addresses
- Dijkstra's algorithm (non-negative weights). Bellman-Ford algorithm (allows negative weights)

ALL PAIRS SHORTEST PATHS

- Find shortest paths between every pair of vertices i & j
- optimal airline, railway, road routes between cities
- Run Dijkstra or Bellman-Ford from each vertex
- Is there another way?

TRANSITIVE CLOSURE

- Adjacency matrix A represents paths of length 1
- Matrix multiplication, $A^2 = A \times A$
 - $A^2[i, j] = 1$, if there is a path of length 2 from i to j
 - For some k , $A[i, k] = A[k, j] = 1$
- In general, $A^{l+1} = A^l \times A$
 - $A^{l+1}[i, j] = 1$, if there is a path of length $l+1$ from i to j
 - For some k , $A^l[i, k] = 1, A[k, j] = 1$
- $A^t = A + A^2 + \dots + A^{n-1}$

AN ALTERNATIVE APPROACH

- $B^k[i, j] = 1$, if there is path from i to j via vertices $\{0, 1, \dots, k-1\}$
 - Constraint applies only to intermediate vertices between i and j
 - $B^0[i, j] = 1$ if there is a direct edge
 - $B^0 = A$
- $B^{k+1}[i, j] = 1$ if
 - $B^k[i, j] = 1$ — can already reach j from i via

$\{0, 1, \dots, k-1\}$

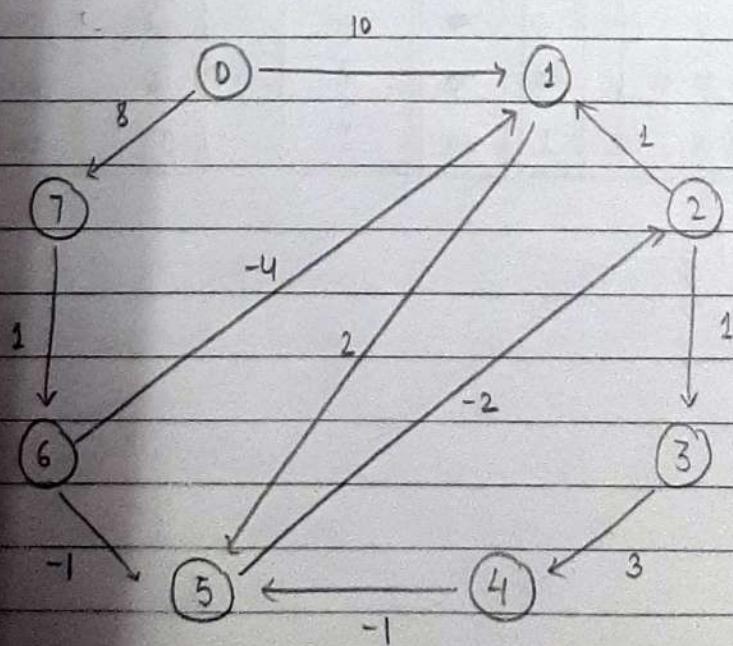
- $B^k[i, k] = 1$ and $B^k[k, j] = 1$ - use $\{0, 1, \dots, k-1\}$ to go from i to k and then from k to j

WARSHALL's ALGORITHM

- $B^k[i, j] = 1$. if there is path from i to j via vertices $\{0, 1, \dots, k-1\}$
- $B^0[i, j] = A[i, j]$
 - Direct edges , no intermediate vertices
- $B^{k+1}[i, j] = 1$ if :
 - $B^k[i, j] = 1$, or
 - $B^k[i, k] = 1$ and $B^k[k, j] = 1$
- The algorithm on the left also computes transitive closure - Warshall's Algorithm
- $B^n[i, j] = 1$. if there is some path from i to j with intermediate vertices in $\{0, 1, \dots, n-1\}$
- $B^n = A^+$
- We adapt Warshall's algorithm to compute all-pairs shortest paths

FLOYD-WARSHALL ALGORITHM

- Let $SP^k[i, j]$ be the length of the shortest path from i to j via vertices $\{0, 1, \dots, k-1\}$
- $SP^0[i, j] = w[i, j]$
 - No intermediate vertices, shortest path is weight of direct edge
 - Assume $w[i, j] = \infty$, if $(i, j) \notin E$
- $SP^{k+1}[i, j]$ is the minimum of
 - $SP^k[i, j] \rightarrow$ shortest path using only $\{0, 1, \dots, k-1\}$
 - $SP^k[i, k] + SP^k[k, j] \rightarrow$ combine shortest path from i to k & k to j
- $SP^n[i, j] = 1$ is the length of the shortest path overall from i to j
 - Intermediate vertices lie in $\{0, 1, \dots, n-1\}$



SP ⁰	0	1	2	3	4	5	6	7	SP ²	0	1	2	3	4	5	6	7	SP ⁷	0	1	2	3	4	5	6	7
0	00	10	00	00	00	00	00	00	0	00	10	00	00	00	12	00	00	0	00	10	10	11	14	12	00	8
1	00	00	00	00	00	00	00	00	1	00	00	00	00	00	2	00	00	1	00	1	0	1	4	2	00	00
2	00	1	00	1	00	00	00	00	2	00	1	00	1	00	3	00	00	2	00	1	1	1	4	3	00	00
3	00	00	00	00	3	00	00	00	3	00	00	00	00	00	3	00	00	3	00	1	0	1	3	2	00	00
4	00	00	00	00	00	-1	00	00	4	00	00	00	00	00	-1	00	00	4	00	-2	-3	-2	1	-1	00	00
5	00	00	-2	00	00	00	00	00	5	00	00	-2	00	00	00	00	00	5	00	-1	-2	-1	2	1	00	00
6	00	-4	00	00	00	-1	00	00	6	00	-4	00	00	00	-2	00	00	6	00	-4	-4	-3	0	-2	00	00
7	00	00	00	00	00	00	1	00	7	00	00	00	00	00	00	1	00	7	00	-3	-3	-2	2	-1	1	00

SP ²	0	1	2	3	4	5	6	7	SP ³	0	1	2	3	4	5	6	7	SP ⁸	0	1	2	3	4	5	6	7
0	00	10	00	00	00	00	00	00	0	00	10	00	00	00	12	00	00	0	00	5	5	6	9	7	9	8
1	00	00	00	00	00	00	00	00	1	00	00	00	00	00	2	00	00	1	00	1	0	1	4	2	00	00
2	00	1	00	1	00	00	00	00	2	00	1	00	1	00	3	00	00	2	00	1	1	1	4	3	00	00
3	00	00	00	00	3	00	00	00	3	00	00	00	00	00	3	00	00	3	00	1	0	1	3	2	00	00
4	00	00	00	00	00	-1	00	00	4	00	00	00	00	00	-1	00	00	4	00	-2	-3	-2	1	-1	00	00
5	00	00	-2	00	00	00	00	00	5	00	-1	-2	-1	00	1	00	00	5	00	-1	-2	-1	0	1	00	00
6	00	-4	00	00	00	-1	00	00	6	00	-4	00	00	00	-2	00	00	6	00	-4	-4	-3	0	-2	00	00
7	00	00	00	00	00	00	1	00	7	00	00	00	00	00	00	1	00	7	00	-3	-3	-2	1	-1	1	00

IMPLEMENTATION

- Shortest path matrix SP is $n \times n \times (n+1)$
- Initialize $SP[i, j, 0]$ to edge weight $w(i, j)$ or ∞ if no edge
- Update $SP[i, j, k]$ from $SP[i, j, k-1]$ using the Floyd Warshall update rule
- Time complexity is $O(n^3)$
- We only need $SP[i, j, k-1]$ to compute $SP[i, j, k]$
- Maintain two "slices" $SP[i, j]$, $SP'[i, j]$, compute SP' from SP , copy SP' to SP , save space

CODE : def floydwarshall (wmat) :

 rows, cols, x = wmat.shape

 infinity = np. max (wmat) * rows * cols + 1

 SP = np. zeros (shape = (rows, cols, cols + 1))

 for i in range (rows) :

 for j in range (cols) :

 SP [i, j, 0] = infinity

 for i in range (rows) :

 for j in range (cols) :

 if wmat [i, j, 0] == type 1

 SP [i, j, 0] = wmat [i, j, 1]

```

for k in range(1, cols+1):
    for i in range(rows):
        for j in range(cols):
            sp[i, j, k] = min(sp[i, j, k-1],
                               sp[i, k-1, k-1] + sp[k-1, j, k-1])
return (sp[:, :, cols])

```

→ SUMMARY ←

- Warshall's algorithm is an alternative way to compute transitive closure
 - $B^k[i, j] = 1$ if we can reach j from i using vertices in $\{0, 1, \dots, k-1\}$
- Adapt Warshall's algorithm to compute all pairs shortest paths
 - $SP^k[i, j]$ is the length of the shortest path from i to j using vertices in $\{0, 1, \dots, k-1\}$
 - $SP^n[i, j]$ is the length of the overall shortest path
 - Floyd-Warshall algorithm
- Works with negative edge weights, assuming no negative cycles
- Simple nested loop implementation, time $O(n^3)$
- Space can be limited to $O(n^2)$ by reusing 2 "slices" SP and SP' .

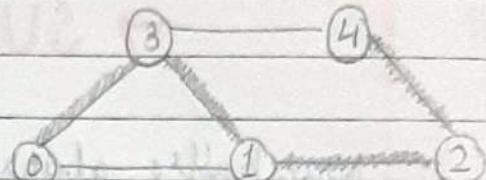
Date
January 27, 2022

Santosh
PAGE NO.
DATE:

Minimum Cost Spanning Trees

SPANNING TREES

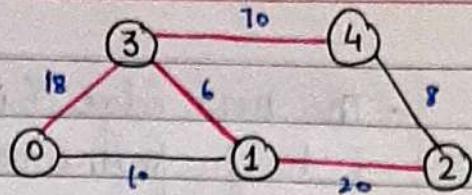
- Retain a minimal set of edges so that graph remains connected
- Recall that a minimally connected graph is a 'tree'.
 - Adding an edge to a tree creates a loop
 - Removing an edge disconnects the graph
- Want a tree that connects all the vertices - spanning tree
- More than one spanning tree , in general



SPANNING TREES WITH COSTS

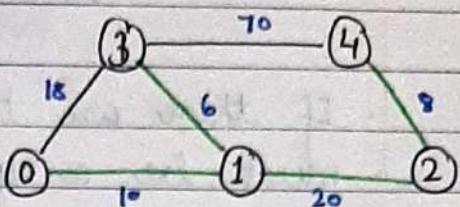
- Restoring a road or laying a fibre optic cable has a cost
- Minimum Cost Spanning Tree
 - Add the cost of all edges in the tree
 - Among the different spanning trees , choose one with

minimum cost.



→ Example

- Spanning tree, Cost is 114 - not minimum cost spanning tree
- Another spanning tree,
Cost is 44, - minimum cost spanning tree



SOME FACTS ABOUT TREES

Definition : A tree is a connected acyclic graph

FACT 1 - A tree on n vertices has exactly $n-1$ edges

- Initially, one single component
- Deleting edge (i, j) must split component (otherwise there is still path from i to j , combine with (i, j) to form cycle)
- Each edge deletion creates one more component
- Deleting $n-1$ edges creates n components, each an isolated vertex

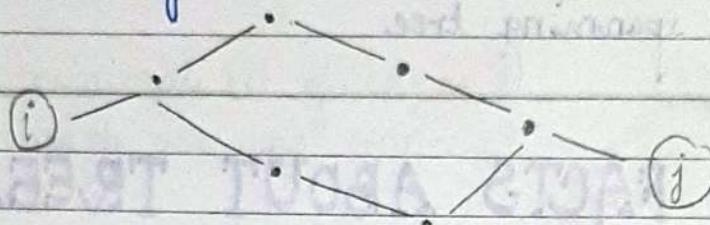
FACT 2 - Adding an edge to a tree must create a cycle.

- Suppose we add an edge (i, j)
- Tree is connected, so there is already a path from i to j

- The new edge (i, j) combined with this path from i to j forms a cycle

FACT 3 - In a tree, every pair of vertices is connected by a unique path.

- If there are two paths from i and to j ; there must be a cycle



Observation : Any 2 of the following facts about a graph G implies the third

- $\rightarrow G$ is connected
- $\rightarrow G$ is acyclic
- $\rightarrow G$ has $n-1$ edges

BUILDING MIN. COST SPANNING TREES

- We will use these facts about trees to build minimum cost spanning trees.
- Two natural strategies
- Start with the smallest edge & 'grow' a tree
 - Prim's Algorithm

- Scan the edges in ascending order of weight to connect components without forming cycles
 - Kruskal's Algorithm

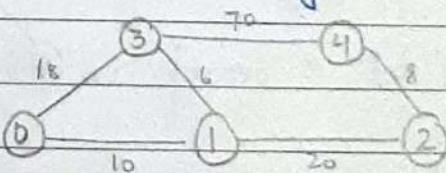
PRIM's ALGORITHM

Minimum Cost Spanning Tree (MCST)

- Weighted undirected graph, $G = (V, E)$, $w: E \rightarrow \mathbb{R}$
 G assumed to be connected
 - Find a minimum cost spanning tree
 - Tree connecting all vertices in V
- STRATEGY 1
 - Incrementally grow the min. cost spanning tree
 - Start with a smallest weight edge overall
 - Extend the current tree by adding the smallest edge from the tree to a vertex not yet in the tree

• EXAMPLE

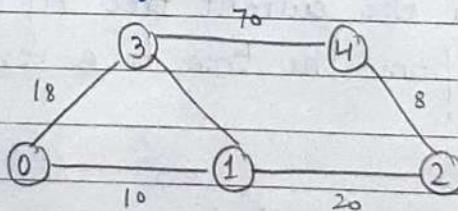
- Start with smallest edge $(1, 3)$
- Extend the tree with $(1, 0)$
- Can't add $(0, 3)$, forms a cycle
- Instead, extend the tree with $(1, 2)$
- Extend the tree with $(2, 4)$.



PRIM's ALGORITHM

- $G = (V, E)$, $W: E \rightarrow \mathbb{R}$
- Incrementally build an MCST
 - $TV \subseteq V$: tree vertices, already added to MCST
 - $TE \subseteq E$: tree edges, already added to MCST
- Initially, $TV = TE = \emptyset$
- Choose minimum weight edge $e = (i, j)$
 - set $TV = \{i, j\}$
 - $TE = \{e\}$ MCST
- Repeat $n-2$ times
 - choose min. weight edge $f = (u, v)$ such that $u \in TV, v \notin TV$
 - Add v to TV , f to TE

EXAMPLE :



$$TV = \{1, 3, 0, 2, 4\}$$

$$TE = \{(1,3), (1,0), (1,2), (2,4)\}$$

CORRECTNESS OF PRIM's ALGORITHM

⇒ Minimum Separator Lemma :

- Let V be partitioned into two non-empty sets U & $W = V \setminus U$
- Let $e = (u, w)$ be the minimum cost edge with $u \in U, w \in W$
- Every MCST must include e

→ Assume for now, all edge weights distinct

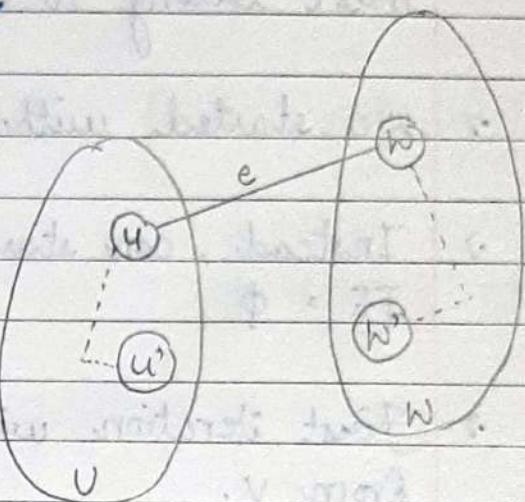
→ Let T be an MCST, $e \notin T$

→ T contains a path p from u to w

• p starts in U , ends in W

• Let $f = (u', w')$ be the first edge on p crossing from U to W .

• Drop f , add e to get a cheaper spanning tree



→ What if two edges have the same weight?

→ Assign each edge a unique index from 0 to $m-1$

→ Define $(e, i) < (f, j)$ if $w(e) < w(f)$ or $w(e) = w(f)$ and $i < j$.

→ In Prim's algorithm, TV and $W = V \setminus TV$ partition V

→ Algorithm picks smallest edge connecting TV & W , which must belong to every MCST

- In fact, for any $v \in V$, $\{v\}$ and $V \setminus \{v\}$ form a partition
- The smallest weight edge leaving any vertex must belong to every MST
- We started with overall minimum cost edge
- Instead, can start at any vertex v , with $TV = \{v\}$
 $TE = \emptyset$
- First iteration will pick minimum cost edge from v .

IMPLEMENTATION

- Keep track of :
 - visited[v] - is v in the spanning tree?
 - distance[v] - shortest dist. from v to tree
 - Tree Edges - edges in the current spanning tree
- Initialize visited[v] to False, distance[v] to infinity
- First add vertex 0 to tree
- Find edge (u, v) leaving the tree where distance[v] is minimum, add it to the tree; update distance[w] of neighbours

CODE :

```
def primlist (WList) :
```

infinity = $1 + \max([d \text{ for } u \text{ in } WList.keys() \text{ for } (v, d) \text{ in } WList[u]])$
 $(visited, distance, TreeEdges) = (1, \{ \}, [])$

for v in $WList.keys()$:

$(visited[v], distance[v]) = (False, infinity)$

$visited[0] = True$

for (v, d) in $WList[0]$:

$distance[v] = d$

for i in $WList.keys()$:

$(mindist, nextv) = (\text{infinity}, \text{None})$

for u in $WList.keys()$:

for (v, d) in $WList[u]$:

if $visited[u]$ and $(\text{not visited}[v])$ and
 $d < mindist$:

$(mindist, nextv, nexte) = (d, v, (u, v))$

if $nextv$ is None :

break

$visited[nextv] = True$

$TreeEdges.append(nexte)$

for (v, d) in $WList[nextv]$:

if $\text{not visited}[v]$:

$distance[v] = \min(distance[v], d)$

return $(TreeEdges)$

COMPLEXITY

- Initialization takes $O(n)$
- Loop to add nodes to the tree runs $O(n)$ times
- Each iteration takes $O(m)$ time to find a node to add
- Overall time is $O(mn)$, which could be $O(n^3)$!
- Can we do better?

IMPROVED IMPLEMENTATION

- For each v , keep track of its nearest neighbour in the tree
 - visited [v] - is v in the spanning tree?
 - distance [v] - shortest distance from v to the tree
 - nbr [v] - nearest neighbour of v in tree.
- Scan all non-tree vertices to find nextv with minimum distance
- Then $(\text{nbr}[\text{nextv}], \text{nextv})$ is the tree edge to add
- Update distance [v] and nbr [v] for all neighbours of nextv.

COMPLEXITY : → Now the scan to find next vertex to add is $O(n)$

- Very similar to Dijkstra's algorithm, except for the update rule for distance

- Like Dijkstra's algorithm, this is still $O(n^2)$ even for adjacency lists
- With a more clever data structure to extract the minimum, we can do better

CODE :

```

def primlist2 (WList):
    infinity = 1 + max ([d for u in WList.keys()
                         for (v,d) in WList[u] ])
    (visited, distance, nbr) = ( { }, { }, { } )

    for v in WList.keys():
        (visited[v], distance[v], nbr[v]) = (False, infinity, -1)

    visited[0] = True
    for (v,d) in WList[0]:
        (distance[v], nbr[v]) = (d, 0)

    for i in range (1, len (WList.keys())):
        nextd = min ([distance[v] for v in WList.keys()
                      if not visited[v] ])
        nextvlist = [v for v in WList.keys()
                     if (not visited[v]) and
                        distance[v] == nextd]

        if nextvlist == []:
            break

        nextv = min (nextvlist)

```

visited[nextv] = True

for (v, d) in Nlist[nextv]:

if not visited[v]:

$(\text{distance}[v], \text{nbr}[v]) = (\min(\text{distance}[v], d),$
nextv)

return (nbr)

→ SUMMARY ←

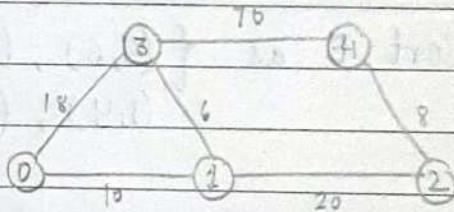
1}

- Prim's algorithm grows an MCST starting with any vertex
- At each step, connect one more vertex to the tree using minimum cost edge from inside the tree to outside the tree
- Correctness follows from Minimum Separator Lemma
- Implementation similar to Dijkstra's algorithms
 - Update rule for distance is different
- Complexity is $O(n^2)$
 - Even with adj. lists
 - Bottleneck is identifying unvisited vertex with min distance.

KRUSKAL's ALGORITHM

Minimum Cost Spanning Tree (MCST)

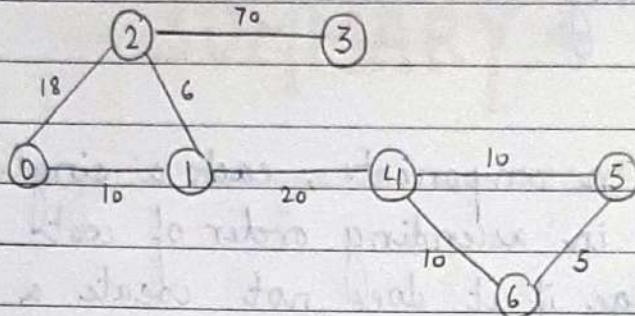
- Weighted undirected graph, $G = (V, E)$, $w: E \rightarrow \mathbb{R}$
 G assumed to be connected
- Find a minimum cost spanning tree
 - Tree connecting all vertices in V
- STRATEGY 2
 - start with n components, each a single vertex
 - Process edges in ascending order of cost
 - Include edge if it does not create a cycle
- EXAMPLE
 - start with smallest edge $(1, 3)$
 - Add next small edge $(2, 4)$
 - Add next smallest edge $(0, 1)$
 - Can't add $(0, 3)$ forms a cycle
 - Add next smallest edge $(1, 2)$



KRUSKAL's ALGORITHM

- $G = (V, E)$, $w: E \rightarrow \mathbb{R}$
- Let $E = \{e_0, e_1, \dots, e_{m-1}\}$ be edges sorted in ascending order by weight

- Let $TE \subseteq E$ be the set of tree edges already added to MCST
- Initially, $TE = \emptyset$
- Scan E from e_0 to e_{m-1}
 - If adding e_i to TE creates a loop, skip it
 - otherwise, add e_i to TE

EXAMPLE

Sort E as $\{(5,6), (1,2), (0,1), (4,5), (4,6), (0,2), (1,4), (2,3)\}$

Set $TE = \emptyset$

Add $(5,6)$, $(1,2)$, $(0,1)$, $(4,5)$; Skip $(4,6)$, $(0,2)$; Add $(1,4)$
 $(2,3)$

$TE = \{(5,6), (1,2), (0,1), (4,5), (1,4), (2,3)\}$

CORRECTNESS OF KRUSKAL'S ALGORITHM

- Minimum separator lemma
 - let V be partitioned into two non-empty sets U and $W = V \setminus U$
 - let $e = (u,w)$ be the minimum cost edge with $u \in U$, $w \in W$

- Every MCST must include e
- Edges in TE partition vertices into connected components
 - Initially each vertex is a separate component
- Adding $e = (u, w)$ merges components of u and w
 - If u and w are in the same component, e forms a cycle & is discarded
- Let U be component of u, W be $V \setminus U$
 - U, W form a partition of V with $u \in U$ & $w \in W$
 - Since we are scanning edges in ascending order of cost, e is minimum cost edge connecting U and W, so it must be part of any MCST.

IMPLEMENTATION

- Collect edges in a list as (d, u, v)
 - Weight as first component for easy sorting
- Main challenge is to keep track of connected components
 - Dictionary to record component of each vertex
 - Initially each vertex is an isolated component
 - When we add an edge (u, v) , merge the components of u & v
- Analysis
 - Sorting the edges is $O(m \log m)$ since m is at most n^2 equivalently, $O(m \log n)$

- Outerloop runs m times
 - Each time we add a tree edge, we have to merge components - $O(n)$ scan
 - $n-1$ tree edges, so this is done $O(n)$ times
- Overall, $O(n^2)$
- Complexity is $O(n^2)$
- Bottleneck is naive strategy to label & merge components
- Components partition vertices
 - collection of disjoint sets
- Data structure to maintain collection of disjoint sets
 - find(v) - return set containing v
 - union(u, v) - merge sets of u, v .
- Efficient union-find brings complexity down to $O(m \log n)$

CODE : def kruskal(WList) :

(edges, component, TE) = ([], {}, [])

for u in WList.keys() :

edges.extend([(d, u, v) for (r, d) in WList[u]])

component[u] = u

edges.sort()

for (d, u, v) in edges :

if component[u] != component[v] :

```

    TE.append ((u,v))
    c = component [u]
    for w in Wlist.keys():
        if component [w] == c:
            component [w] = component [v]
    return (TE)

```

→ SUMMARY ←

- Kruskal's algorithm builds an MCST bottom up
 - start with n components, each an isolated vertex
 - scan edges in ascending order of cost
 - whenever an edge merges disjoint components, add it to the MCST
- correctness follows from Minimum Separator lemma.
- complexity is $O(n^2)$ due to naive handling of components
 - will see how to improve to $O(m \log n)$
- If edge weights repeat, MCST is not unique
- "choose minimum cost edge" will allow choices
 - consider a triangle on 3 vertices with all ^{edges} equal
- Different choices lead to different spanning trees
- In general, there may be a very large number of MCST