


▼ Week 5

▼ SHORTEST PATH IN WEIGHTED GRAPHS

Weighted graphs

- BFS explores a graph level by level
- BFS computes shortest path, in terms of number of edges, to every reachable vertex
- May assign values to edges
 - Cost, time, distance,...
 - Weighted graphs
- $G = (V, E), W : E \rightarrow \mathbb{R}$



- Adjacency matrix: Record the weights along with edge information - weight is always 0 if no edge

	0	1	2	3	4	5	6
0	(0,0)	(1,10)	(1,80)	(0,0)	(0,0)	(0,0)	(0,0)
1	(1,10)	(0,0)	(1,6)	(0,0)	(1,20)	(0,0)	(0,0)
2	(1,80)	(1,6)	(0,0)	(1,70)	(0,0)	(0,0)	(0,0)
3	(0,0)	(0,0)	(1,70)	(0,0)	(0,0)	(0,0)	(0,0)
4	(0,0)	(1,20)	(0,0)	(0,0)	(0,0)	(1,50)	(1,5)
5	(0,0)	(0,0)	(0,0)	(0,0)	(1,50)	(0,0)	(1,10)
6	(0,0)	(0,0)	(0,0)	(0,0)	(1,5)	(1,10)	(0,0)

- Adjacency list: Record weights along with the edge information

0	[(1,10),(2,80)]
1	[(0,10),(2,6),(4,20)]
2	[(0,80),(1,6),(3,70)]
3	[(2,70)]
4	[(1,20),(5,50),(6,5)]
5	[(4,50),(6,10)]
6	[(4,5),(5,10)]

Shortest paths in weighted graphs

- BFS computes shortest path, in terms of number of edges, to every reachable vertex
- In a weighted graph, add up the weights along a path
- Weighted shortest path need not have minimum number of edges
 - Shortest path from 0 to 2 is via 1 (weight = 16)

Shortest path problems

Single source shortest paths

- Find shortest paths from a fixed vertex to every other vertex
- Transport finished product from factory (single source) to all retail outlets
- Courier company delivers items from distribution centre (single source) to addresses

All pairs shortest path

- Find shortest paths between every pair of vertices i and j
- Optimal airline, railway, road routes between cities

Negative edges weights

- Can negative edge weights be meaningful?
- Taxi driver trying to head home at the end of the day
 - Roads with few customer, drive empty (positive weight)
 - Roads with many customers, make profit (negative weight)
 - Find route towards home that minimizes cost

Negative cycles

- A negative cycle is one whose weight is negative
 - Sum of the weights of edges that make up the cycle
- By repeatedly traversing a negative cycle, total cost keeps decreasing
- If a graph has a negative cycle, total cost keeps decreasing

- Without negative cycles, we can compute shortest paths even if some weights are


Summary

- In a weighted graph, each edge has a cost
 - Entries in adjacency matrix capture edge weights
- Length of a path is the sum of the weights
 - Shortest path in a weighted graph need not be minimum in terms of number of edges
- Different shortest path problems
 - Single source - from one designated vertex to all others
 - All-pairs - Between every pair of vertices
- Negative edge weights
 - Should not have negative cycles
 - Without negative cycles, shortest paths still well defined


▼ SINGLE SOURCE SHORTEST PATHS

Single source shortest paths

- Weighted graph:
 - $G = (V, E)$
 - $W : E \rightarrow \mathbb{R}$
- Single source shortest paths
 - Find shortest paths from a fixed vertex to every other vertex
- Assume, for now, that edge weights are all non-negative
- Compute shortest path from 0 to all other vertices
- Imagine vertices are oil depots, edges are pipelines
- Set fire to oil depot at vertex 0
- Fire travels at uniform speed along each pipeline
- First oil depot to catch fire after 0 is nearest vertex
- Next oil depot is second nearest vertex
- ...




- Compute expected burn time for each vertex
- Each time a new vertex burns, update the expected burn times of its neighbours
- Algorithm due to Edsger W Dijkstra



Dijkstra's algorithm: Proof of correctness

- Each new shortest path we discover extends an earlier one (Greedy method)
- By induction, assume we have found shortest paths to all vertices already burnt
- Next vertex to burn is v , via x
- Can't find a shorter path later from y to v via w
 - Burn time of $w \geq$ burn time of v
 - Edge from w to v has weight ≥ 0
- This argument breaks down if edge (w, v) can have negative weight
 - Can't use Dijkstra's algorithm with negative edge weights



Implementation

- Maintain 2 dictionaries with vertices as keys
 - `visited` initially `False` for all v (burnt vertices)
 - `distance` initially `infinity` for all v (expected burn time)
- Set `distance[s]` to 0
- Repeat, until all reachable vertices are visited
 - Find unvisited vertex `nextv` with minimum distance
 - Set `visited[nextv]` to `True`
 - Recompute `distance[v]` for every neighbour v of `nextv`

```
def dijkstra(WMat, s):
    # s is the source vertex; WMat is the weighted adj matrix
    (rows,cols,x) = WMat.shape # x is the edge/weight info
    # x[0] edge info; x[1] weight info
    infinity = np.max(WMat) * rows + 1
    # max value in the matrix multiplied by rows + 1 is larger than all
    (visited, distance) = ({},{})
    for v in range(rows):
        (visited[v], distance[v]) = (False, infinity)
    distance[s] = 0
    for u in range(rows):
        nextd = min([distance[v] for v in range(rows) if not visited[v]])
        nextvlist = [v for v in range(rows) if(not visited[v]) and distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for v in range(cols):
            if WMat[nextv,v,0] == 1 and (not visited[v]):
```

```

distance[v] = min(distance[v], distance[nextv] + WMat[nextv, v, 1])
return(distance)

```

Complexity

- Setting `infinity` takes $O(n^2)$ time
- Main loop runs n times
 - Each iteration visits one more vertex
 - $O(n)$ to find next vertex to visit
 - $O(n)$ to update `distance[v]` for neighbours
- Overall $O(n^2)$
- If we use an adjacency list
 - Setting `infinity` and updating distances both $O(m)$, amortised
 - $O(n)$ bottleneck remains to find next vertex to visit
 - Better data structure? Later

```

def dijkstralist(WList, s):
    infinity = 1 + len(WList.keys()) * max([d for u in WList.keys() for (v,d) in WList[u]])
    (visited,distance) = ({},{})
    for v in WList.keys():
        (visited[v], distance[v]) = (False,infinity)
    distance[s] = 0
    for u in WList.keys():
        nextd = min([distance[v] for v in WList.keys() if not visited[v]])
        nextvlist = [v for v in WList.keys() if (not visited[v]) and distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for (v,d) in WList[nextv]:
            if not visited[v]:
                distance[v] = min(distance[v], distance[nextv])
    return(distance)

```

Summary

- Dijkstra's algorithm computes single source shortest paths
- Use a greedy strategy to identify vertices to visit
 - Next vertex to visit is based on shortest distance computed so far
 - Need to prove that such a strategy is correct
 - Correctness requires edge weights to be non-negative
- Complexity is $O(n^2)$
 - Even with adjacency lists
 - Bottleneck is identifying unvisited vertex with minimum distance
 - Need a better data structure to identify and remove minimum (or maximum) from a collection

▼ SINGLE SOURCE SHORTEST PATHS WITH NEGATIVE WEIGHTS

Dijkstra's Algorithm

Burning pipe analogy

- We keep track of the following
 - The vertices that have been burnt
 - The expected burn time of vertices
- Initially
 - No vertex is burnt
 - Expected burn time of vehicles
 - Expected burn time of rest is ∞
- While there are vertices yet to burn
 - Pick unburnt vertex with minimum expected burn time, mark it s burnt
 - Update the expected burn time of its neighbours

Initialization (assume source vertex 0)

- $B(i) = \text{False}$, for $0 \leq i < n$
 - $UB = \{k \mid B(k) = \text{False}\}$
- $EBT(i) = \begin{cases} 0, & \text{if } i = 0 \\ \infty & \text{otherwise.} \end{cases}$

Update, if $UB \neq \emptyset$

- Let $j \in UB$ such that $EBT(j) \leq EBT(k)$ for all $k \in UB$
- Update $B(j) = \text{True}, UB = UB \setminus \{j\}$
- For each $(j, k) \in E$ such that $k \in UB$,
 $EBT(k) = \min(EBT(k), EBT(j) + W(j, k))$

Extending to negative edge weights

- The difficulty with negative edge weights is that we stop updating the burn time once a vertex is burnt
- What if we allow updates even after a vertex is burnt?
- Negative edge weights are allowed, but no negative cycles
- Going around a cycle can only add to the length
- Shortest route to every vertex is a path, no loops
- Suppose minimum weight path from 0 to k is

$$0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1} \xrightarrow{w_{\ell}} k$$

- Need not be minimum in terms of number of edges
- Every prefix of this path must be a minimum weight path

■ $0 \xrightarrow{w_1} j_1$
 ■ $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2$
 ■ \dots
 ■ $0 \xrightarrow{w_1} j_1 \xrightarrow{w_2} j_2 \xrightarrow{w_3} \dots \xrightarrow{w_{\ell-1}} j_{\ell-1}$

- Once we discover shortest path to j_{l-1} , next update will fix shortest path to k
- Repeatedly update shortest distance to each vertex based on shortest distance to its neighbours
 - Update cannot push this distance below actual shortest distance
- After l updates, all shortest paths using $\geq l$ edges have stabilized
 - Minimum weight path to any node has at most $n - 1$ edges
 - After $n - 1$ update, all shortest paths have stabilized

Bellman-Ford Algorithm

Initialization (source vertex 0)

- $D(j)$: minimum distance known so far to vertex j
- $D(j) = \begin{cases} 0, & \text{if } j = 0 \\ \infty & \text{otherwise.} \end{cases}$

Repeat $n - 1$ times

- For each vertex $j \in \{0, 1, \dots, n - 1\}$, for each edge $(j, k) \in E$,
 $D(k) = \min(D(k), D(j) + W(j, k))$

Works for directed and undirected graphs


```
def bellmanford(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows + 1
    distance = {}
    for v in range(rows):
        distance[v] = infinity
    distance[s] = 0
    for i in range(rows):
```



```


for u in range(rows):
    for v in range(cols):
        if WMat[u,v,0] == 1:
            distance[v] = min(distance[v], distance[u] + WMat[u,v,1])
return(distance)

```



v	$D(v)$							
0	0							
1	∞							
2	∞							
3	∞							
4	∞							
5	∞							
6	∞							
7	∞							


■ Initialize $D(0) = 0$



v	$D(v)$							
0	0	0	0					
1	∞	10	10					
2	∞	∞	∞					
3	∞	∞	∞					
4	∞	∞	∞					
5	∞	∞	12					
6	∞	∞	9					
7	∞	8	8					

- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update

$$D(k) = \min(D(k), D(j) + W(j, k))$$




Navigation icons: back, forward, search, etc.

v	$D(v)$							
0	0	0	0	0				
1	∞	10	10	5				
2	∞	∞	∞	10				
3	∞	∞	∞	∞				
4	∞	∞	∞	∞				
5	∞	∞	12	8				
6	∞	∞	9	9				
7	∞	8	8	8				


- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update

$$D(k) = \min(D(k), D(j) + W(j, k))$$



Navigation icons: back, forward, search, etc.

v	$D(v)$							
0	0	0	0	0	0	0	0	0
1	∞	10	10	5	5	5	5	5
2	∞	∞	∞	10	6	5	5	5
3	∞	∞	∞	∞	11	7	6	6
4	∞	∞	∞	∞	∞	14	10	9
5	∞	∞	12	8	7	7	7	7
6	∞	∞	9	9	9	9	9	9
7	∞	8	8	8	8	8	8	8



- Initialize $D(0) = 0$
- For each $(j, k) \in E$, update

$$D(k) = \min(D(k), D(j) + W(j, k))$$

- What if there was a negative cycle? Distance would continue to decrease
- Check if update n reduces any $D(v)$

Complexity

- Initialising `infinity` takes $O(n^2)$ time
- The outer update loop runs $O(n)$ times
- In each iteration, we have to examine every edge in the graph
 - This take $O(n^2)$ for an adjacency matrix
- Overall, $O(n^3)$
- If we shift to adjacency lists
 - Initializing `infinity` is $O(m)$
 - Scanning all edges in each update iteration is $O(m)$

```
def bellmanfordlist(WList,s):
    infinity = 1 + len(WList.keys())*max([d for u in WList.keys() for (v,d) in WList[u]])
    distance = {}
    for v in WList.keys():
        distance[v] = infinity

    distance[s] = 0

    for i in WList.keys():
        for u in WList.keys():
            for (v,d) in WList[u]:
```

```
distance[v] = min(distance[v], distance[u] + d)
return(distance)
```

Summary

- Dijkstra's algorithm assumes non-negative edge weights
 - Final distance is frozen each time a vertex "burns"
- Without negative cycles, every shortest route is a path
- Every prefix of a shortest path is also a shortest path
- Iteratively find shortest paths of length $1, 2, \dots, n - 1$
- Update distance to each vertex with every iteration - *Bellman-Ford algorithm*
- $O(n^3)$ time with adjacency matrix, $O(mn)$ time with adjacency list
- If Bellman-Ford algorithm does not converge after $n - 1$ iterations, there is a negative cycle

▼ ALL PAIRS SHORTEST PATHS (FLOYD-WARSHALL ALGORITHM)

Shortest paths in weighted graphs

Two types of shortest path problems of interest

Single source shortest paths

- Find shortest paths from a fixed vertex to every other vertex
- Transport finished product from factory (single source) to all retail outlets
- Courier company delivers items from distribution centre (single source) to addresses

All pairs shortest paths

- Find shortest paths between every pair of vertices i and j
- Optimal airline, railway, road routes between cities
- Run Dijkstra or Bellman-Ford from each vertex
- Is there is another way?

Transitive closure

- Adjacency matrix A represents paths of length 1
- Matrix multiplication, $A^2 = A \times A$
 - $A^2[i, j] = 1$ if there is a path of length 2 from i to j
 - For some k , $A[i, k] = 1$, $A[k, j] = 1$
- $A^+ = A + A^2 + \dots + A^{n-1}$

An alternative approach

- $B^k[i, j] = 1$ if there is path from i to j via vertices $\{0, 1, \dots, k - 1\}$

- Constraint applies only to intermediate vertices between i and j
- $B^0[i, j] = 1$ if there is a direct edge
- $B^0 = A$
- $B^{k+1}[i, j] = 1$ if
 - $B^k[i, j] = 1$ - can already reach j from i via $\{0, 1, \dots, k-1\}$
 - $B^k[i, k] = 1$ and $B^k[k, j] = 1$ - use $\{0, 1, \dots, k-1\}$ to go from i to k and then from k to j


Warshall's Algorithm

- $B^k[i, j] = 1$ if there is path from i to j via vertices $\{0, 1, \dots, k-1\}$
- $B^0[i, j] = A[i, j]$
 - Direct edges, no intermediate vertices
- $B^{k+1}[i, j] = 1$ if
 - $B^k[i, j] = 1$, or
 - $B^k[i, k] = 1$ and $B^k[k, j] = 1$
- This algorithm also computes transitive closure - Warshall's algorithm
- $B^n[i, j] = 1$ if there is some path from i to j with intermediate vertices in $\{0, 1, \dots, n-1\}$
- $B^n = A^+$
- We adapt Warshall's algorithm to compute all-pairs shortest paths

Floyd-Warshall Algorithm

- Let $SP^k[i, j]$ be the length of the shortest path i to j via vertices $\{0, 1, \dots, k-1\}$
- $SP^0[i, j] = W[i, j]$
 - No intermediate vertices, shortest path is weight of direct edge
 - Assume $W[i, j] = \infty$ if $(i, j) \notin E$
- $SP^{k+1}[i, j]$ is the minimum of
 - $SP^k[i, j]$
Shortest path using only $\{0, 1, \dots, k-1\}$
 - $SP^k[i, k] + SP^k[k, j]$
Combine shortest path from i to k and k to j

SP^0	0	1	2	3	4	5	6	7
0	∞	10	∞	∞	∞	∞	∞	8
1	∞	∞	∞	∞	∞	2	∞	∞
2	∞	1	∞	1	∞	∞	∞	∞
3	∞	∞	∞	∞	3	∞	∞	∞
4	∞	∞	∞	∞	∞	-1	∞	∞
5	∞	∞	-2	∞	∞	∞	∞	∞
6	∞	-4	∞	∞	∞	-1	∞	∞
7	∞	∞	∞	∞	∞	∞	1	∞



SP^1	0	1	2	3	4	5	6	7
0	∞	10	∞	∞	∞	∞	∞	8
1	∞	∞	∞	∞	∞	2	∞	∞
2	∞	1	∞	1	∞	∞	∞	∞
3	∞	∞	∞	∞	3	∞	∞	∞
4	∞	∞	∞	∞	∞	-1	∞	∞
5	∞	∞	-2	∞	∞	∞	∞	∞
6	∞	-4	∞	∞	∞	-1	∞	∞
7	∞	∞	∞	∞	∞	∞	1	∞

SP^2	0	1	2	3	4	5	6	7
0	∞	10	∞	∞	∞	12	∞	8
1	∞	∞	∞	∞	∞	2	∞	∞
2	∞	1	∞	1	∞	3	∞	∞
3	∞	∞	∞	∞	3	∞	∞	∞
4	∞	∞	∞	∞	∞	-1	∞	∞
5	∞	∞	-2	∞	∞	∞	∞	∞
6	∞	-4	∞	∞	∞	-2	∞	∞
7	∞	∞	∞	∞	∞	∞	1	∞

SP^3	0	1	2	3	4	5	6	7
0	∞	10	∞	∞	∞	12	∞	8
1	∞	∞	∞	∞	∞	2	∞	∞
2	∞	1	∞	1	∞	3	∞	∞
3	∞	∞	∞	∞	3	∞	∞	∞
4	∞	∞	∞	∞	∞	-1	∞	∞
5	∞	-1	-2	-1	∞	1	∞	∞
6	∞	-4	∞	∞	∞	-2	∞	∞
7	∞	∞	∞	∞	∞	∞	1	∞

SP^7	0	1	2	3	4	5	6	7
0	∞	10	10	11	14	12	∞	8
1	∞	1	0	1	4	2	∞	∞
2	∞	1	1	1	4	3	∞	∞
3	∞	1	0	1	3	2	∞	∞
4	∞	-2	-3	-2	1	-1	∞	∞
5	∞	-1	-2	-1	2	1	∞	∞
6	∞	-4	-4	-3	0	-2	∞	∞
7	∞	-3	-3	-2	1	-1	1	∞

SP^8	0	1	2	3	4	5	6	7
0	∞	5	5	6	9	7	9	8
1	∞	1	0	1	4	2	∞	∞
2	∞	1	1	1	4	3	∞	∞
3	∞	1	0	1	3	2	∞	∞
4	∞	-2	-3	-2	1	-1	∞	∞
5	∞	-1	-2	-1	2	1	∞	∞
6	∞	-4	-4	-3	0	-2	∞	∞
7	∞	-3	-3	-2	1	-1	1	∞

Implementation

- Shortest path matrix SP is $n \times n \times (n + 1)$
- Initialize $SP[i, j, 0]$ to edge weight $W(i, j)$, or ∞ if no edge
- Update $SP[i, j, k]$ from $SP[i, j, k - 1]$ using the Floyd-Warshall update rule
- Time complexity is $O(n^3)$
- We only need $SP[i, j, k - 1]$ to compute $SP[i, j, k]$
- Maintain two "slices" $SP[i, j]$, $SP'[i, j]$, compute SP' from SP , copy SP' to SP , save space

```
def floydwarshall(WMat):
    (rows, cols, x) = WMat.shape
    infinity = np.max(WMat) * rows * rows + 1
    SP = np.zeros(shape=(rows, cols, cols+1))
    for i in range(rows):
        for j in range(cols):
            SP[i, j, 0] = infinity
    for i in range(rows):
        for j in range(cols):
            if WMat[i, j, 0] == 1:
                SP[i, j, 0] = WMat[i, j, 1]

    for k in range(1, cols+1):
        for i in range(rows):
            for j in range(cols):
                SP[i, j, k] = min(SP[i, j, k-1], SP[i, k-1, k-1] + SP[k-1, j, k-1])
    return(SP[:, :, cols])
```


Summary

- Warshall's algorithm is an alternative way to compute transitive closure
 - $B^k[i, j] = 1$ if we can reach j from i using vertices in $\{0, 1, \dots, k-1\}$
- Adapt Warshall's algorithm to compute all pairs shortest paths
 - $SP^k[i, j]$ is the length of the shortest path from i to j using vertices in $\{0, 1, \dots, k-1\}$
 - $SP^n[i, j]$ is the length of the overall shortest path
 - Floyd-Warshall algorithm
- Works with negative edge weights, assuming no negative cycles
- Simple nested loop implementation, time $O(n^3)$
- Space can be limited to $O(n^2)$ by reusing two "slices" SP and SP'

▼ MINIMUM COST SPANNING TREES


Spanning trees

- Retain a minimal set of edges so that graph remains connected
- A minimally connected graph is a tree
 - Adding an edge to a tree creates a loop
 - Removing an edge disconnects the graph
- Want a tree that connects all the vertices - spanning tree



Spanning trees with costs

- Restoring a road or laying a fibre optic cable has a cost
- Minimum cost spanning tree
 - Add the cost of all edges in the tree
 - Among the different spanning trees, choose one with minimum cost
- Example
 - Spanning tree, Cost is 114 - not minimum cost spanning tree
 - Another spanning tree, Cost is 44 - minimum cost spanning tree



Some facts about trees

Defn: A tree is a connected acyclic graph

Fact 1: A tree on n vertices has exactly $n - 1$ edges


- Initially, one single component
- Deleting edge (i, j) must split component
 - Otherwise, there is still a path from i to j , combine with (i, j) to form cycle
- Each edge deletion creates one more component
- Deleting $n - 1$ edges creates n components, each an isolated vertex

Fact 2: Adding an edge to a tree must create a cycle

- Suppose we add an edge (i, j)
- Tree is connected, so there is already a path from i to j forms a cycle

Fact 3: In a tree, every pair of vertices is connected by a unique path

- If there are two paths from i to j , there must be a cycle



Observation: Any two of the following facts about a graph G implies the third

- G is connected
- G is acyclic
- G has $n - 1$ edges

Summary

- We will use these facts about trees to build minimum cost spanning trees
- Two natural strategies
- Start with the smallest edge and "grow" a tree


- Prim's algorithm
- Scan the edges in ascending order of weight to connect components without forming cycles
- Kruskal's algorithm

▼ MINIMUM COST SPANNING TREES - PRIM'S ALGORITHM

Minimum Cost Spanning tree (MCST)

- Weighted undirected graph, $G = (V, E)$, $W: E \rightarrow \mathbb{R}$
 - G assumed to be connected
- Find a minimum cost spanning tree
 - Tree connecting all vertices in V
- *Strategy*
 - Incrementally grow the minimum cost spanning tree
 - Start with a smallest weight edge overall
 - Extend the current tree by adding the smallest edge from the tree to a vertex not yet in the tree

Example




- Start with smallest edge (1, 3)
- Extend the tree with (1, 0)
- Can't add (0, 3), forms a cycle
- Instead, extend the tree with (1, 2)
- Extend the tree with (2, 4)

Prim's algorithm

- $G = (V, E)$, $W: E \rightarrow \mathbb{R}$
- Incrementally build an MCST
 - $TV \subseteq V$: tree vertices, already added to MCST
 - $TE \subseteq E$: tree edges, already added to MCST

- Initially, $TV = TE = \emptyset$
- Choose minimum weight edge $e = (i, j)$
 - Set $TV = \{i, j\}, TE = \{e\}$ MCST
- Repeat $n - 2$ times
 - Choose minimum weight edge $f = (u, v)$ such that $u \in TV, v \notin TV$
 - Add v to TV, f to TE

Example




$$TV = \{1, 3, 0, 2\}$$

$$TE = \{(1, 3), (1, 0), (1, 2)\}$$

Correctness of Prim's Algorithm

Minimum Separator Lemma

- Let V be partitioned into two non-empty sets U and $W = V \setminus U$
- Let $e = (u, w)$ be the minimum cost edge with $u \in U, w \in W$
- Every MCST must include e
- Assume for now, all edge weights distinct
- Let T be an MCST, $e \notin T$
- T contains a path p from u to
 - p starts in U , ends in W
 - Let $f = (u', w')$ be the first edge on p crossing from U to W
 - Drop f , add e to get a cheaper spanning tree
- What if two edges have same weights?
- Assign each edge a unique index from 0 to $m - 1$
- Define $(e, i) < (f, j)$ if $W(e) < W(f)$ or $W(e) = W(f)$ and $i < j$



- In Prim's algorithm, TV and $W = V \setminus TV$ partition V
- Algorithm picks smallest edge connecting TV and W , which must belong to every MCST
- In fact, for any $v \in V$, $\{v\}$ and $V \setminus \{v\}$ form a partition
- The smallest weight edge leaving any vertex must belong to every MCST
- We started with overall minimum cost edge
- Instead, can start at any vertex v , with $TV = \{v\}$ and $TE = \emptyset$

Implementation

1)

- Keep track of
 - `visited[v]` - is v in the spanning tree?
 - `distance[v]` - shortest distance from v to the tree
 - `TreeEdges` - edges in the current spanning tree

```
def primlist(WList):
    infinity = 1 + max([d for u in WList.keys() for (v,d) in WList[u]])
    (visited, distance, TreeEdges) = ({}, {}, [])
    for v in WList.keys():
        (visited[v], distance[v]) = (False, infinity)
    visited[0] = True
    for (v,d) in WList[0]:
        distance[v] = d
    for i in WList.keys():
        (mindist, nextv) = (infinity, None)
        for u in WList.keys():
            for (v,d) in WList[u]:
                if visited[u] and (not visited[v]) and d < mindist:
                    (mindist, nextv, nexte) = (d, v, (u,v))
    if nextv is None:
        break
```

```

visited[nextv] = True
TreeEdges.append(nexte)
for (v,d) in WList[nextv]:
    if not visited[v]:
        distance[v] = min(distance[v], d)
return(TreeEdges)

```

- Initialize `visited[v]` to `False`, `distance[v]` to `infinity`
- First add vertex 0 to tree
- Find edge (u, v) leaving the tree where `distance[v]` is minimum, add it to the tree, update `distance[w]` of neighbours

Complexity

- Initialization takes $O(n)$
- Loop to add nodes to the tree runs $O(n)$ times
- Each iteration takes $O(m)$ time to find a node to add
- Overall time is $O(mn)$, which could be $O(n^3)$!

2)

- For each v , keep track of its nearest neighbour in the tree
 - `visited[v]` - is v in the spanning tree?
 - `distance[v]` - shortest distance from v to the tree
 - `nbr[v]` - nearest neighbour of v in tree
- Scan all non-tree vertices to find `nextv` with minimum distance
- Very similar to Dijkstra's algorithm, except for the update rule for distance
- Like Dijkstra's algorithm, this is still $O(n^2)$ even for adjacency lists
- With a more clever data structure to extract the minimum, we can do better

```

def primlist2(WList):
    infinity = 1 + max([d for u in WList.keys() for (v,d) in WList[u]])
    (visited, distance, nbr) = ({}, {}, {})
    for v in WList.keys():
        (visited[v], distance[v], nbr[v]) = (False, infinity, -1)
    visited[0] = True
    for (v,d) in WList[0]:
        (distance[v], nbr[v]) = (d, 0)
    for i in range(1, len(WList.keys())):
        nextd = min([distance[v] for v in WList.keys() if not visited[v]])
        nextvlist = [v for v in WList.keys() if (not visited[v]) and distance[v] == nextd]
        if nextvlist == []:
            break
        nextv = min(nextvlist)
        visited[nextv] = True
        for (v,d) in WList[nextv]:

```

```

if not visited[v]:
    (distance[v],nbr[v]) = (min(distance[v], d), nextv)
return(nbr)

```

Summary

- Prim's algorithm grows an MCST starting with any vertex
- At each step, connect one more vertex to the tree using minimum cost edge from inside the tree to outside the tree
- Correctness follows from Minimum Separator Lemma
- Implementation similar to Dijkstra's algorithms
 - Update rule for distance is different
- Complexity is $O(n^2)$
 - Even with adjacency lists
 - Bottleneck is identifying unvisited vertex with minimum distance
 - Need a better data structure to identify and remove minimum (or maximum) from a collection


▼ MINIMUM COST SPANNING TREES - KRUSAL'S ALGORITHM

Minimum Cost Spanning tree (MCST)

Strategy 2

- Start with n components, each a single vertex
- Process edges in ascending order of cost
- Include edge if it does not create a cycle

Example




- Start with smallest edge, (1, 3)
- Add next smallest edge, (2, 4)
- Add next smallest edge, (0, 1)
- Can't add (0, 3), forms a cycle
- Add next smallest edge, (1, 2)

Kruskal's Algorithm

- $G = (V, E), W : E \rightarrow \mathbb{R}$
- Let $E = \{e_0, e_1, \dots, e_{m-1}\}$ be edges sorted in ascending order by weight
- Let $TE \subseteq E$ be the set of tree edges already added to MCST
- Initially, $TS = \emptyset$
- Scan E from e_0 to e_{m-1}
 - If adding e_i to TE creates a loop, skip it
 - Otherwise, add e_i to TE

Example



- Sort E as $\{(5, 6), (1, 2), (0, 1), (4, 5), (0, 2), (1, 4), (2, 3)\}$
- Set $TE = \emptyset$
- Keep adding each into TE after checking
- $TE = \{(5, 6), (1, 2), (0, 1), (4, 5), (1, 4), (2, 3)\}$

Correctness of Kruskal's Algorithm

From *Minimum Separator Lemma*

- Edges in TE partition vertices into connected components
 - Initially each vertex is a separate component
- Adding $e = (u, w)$ merges components of u and w
 - If u and w are in the same component, e forms a cycle and is discarded
- Let U be component of u , W be $V \setminus U$
 - U, W form a partition of V with $u \in U$ and $w \in W$
 - Since we are scanning edges in ascending order of cost, e is minimum cost edge connecting U and W , so it must be part of any MCST

Implementation

- Collect edges in a list as (d, u, v)

- Weight as first component for easy sorting
- Main challenge is to keep track of connected components
 - Dictionary to record component of each vertex
 - Initially each vertex is an isolated component
 - When we add an edge (u, v) , merge the components of u and v

```
def kruskal(WList):
    (edges, component, TE) = ([], {}, [])
    for u in WList.keys():
        # Weight as first component to sort easily
        edges.extend([(d,u,v)] for (v,d) in WList[u])
        component[u] = u
    edges.sort()
    print(edges)

    for (d,u,v) in edges:
        if component[u] != component[v]:
            TE.append((u,v))
            c = component[u]
            for w in WList.keys():
                if component[w] == c:
                    component[w] = component[v]
    return(TE)
```

Analysis

- Sorting the edges is $O(m \log m)$
 - Since m is at most n^2 , equivalently $O(m \log n)$
- Outer loop runs m times
 - Each time we add a tree edge, we have to merge components - $O(n)$ scan
 - $n - 1$ tree edges, so this is done $O(n)$ times
- Overall $O(n^2)$
- Bottleneck is naive strategy to label and merge components
- Components partition vertices
 - Collection of disjoint sets
- Data structure to maintain collection of disjoint sets
 - `find(v)` - return set containing v
 - `union(u,v)` - merge sets of u, v
- Efficient union-find brings complexity down to $O(m \log n)$

Summary

- Kruskal's algorithm builds an MCST bottom up

- Start with n components, each an isolated vertex
- Scan edges in ascending order of cost
- Whenever an edge merges disjoint components, add it to the MCST
- Correctness follows from Minimum Separator Lemma
- Complexity is $O(n^2)$ due to naive handling components
 - Will see how to improve to $O(m \log n)$
- If edge weights repeat, MCST is not unique
- "Choose minimum cost edge" will allow choices
 - Consider a triangle on 3 vertices with all edges equal
- Different choices lead to different spanning trees
- In general, there may be a very number of minimum cost spanning trees