

▼ Union-Find Data Structure

Kruskal's Algorithm for Minimum Cost Spanning Tree (MCST)

- Process the edges in ascending order of cost
- If edge (u, v) does not create a cycle, add it
 - (u, v) can be added if u and v are in different components
 - Adding edge (u, v) merges these components
- How can we keep track of components and merge them efficiently?
- Components **partition** vertices
 - Collection of disjoint sets
- Need data structure to maintain collection of disjoint sets
 - `find(v)` - return set containing `v`
 - `union(u, v)` - merge sets of `u`, `v`

Union-Find Data Structure

- A set S partitioned into components $\{C_1, C_2, \dots, C_k\}$
 - Each $s \in S$ belongs to exactly one C_j
- Support the following operations
 - `MakeUnionFind(S)` - set up initial singleton components $\{s\}$, for each $s \in S$
 - `Find(s)` - return the components containing s
 - `Union(s, s')` - merges components containing s , s'

Naive Implementation

- Assume $S = \{0, 1, \dots, n - 1\}$
- Set up an array/dictionary `Component`
- `MakeUnionFind(S)`
 - Set `Component[i] = i` for each `i`
- `Find(i)`
 - Return `Component[i]`
- `Union(i, j)`

```

c_old = Component[i]
c_new = Component[j]

for k in range(n):
    if Component[k] == c_old:
        Component[k] = c_new

```

Complexity

- `MakeUnionFind(S)` - $O(n)$
- `Find(i)` - $O(1)$
- `Union(i, j)` - $O(n)$

Improved Implementation

- Another array/dictionary `Members`
- For each component c , `Members[c]` is a list of its members
- `Size[c] = length(Members[c])` is the number of members
- `MakeUnionFind(S)`
 - Set `Component[i] = i` for all i
 - Set `Members[i] = i`, `Size[i] = 1` for all i
- `Find(i)`
 - Return `Component[i]`
- `Union(i, j)`

```

c_old = Component[i]
c_new = Component[j]

for k in Members[c_old]:
    Component[k] = c_new
    Members[c_new].append(k)
    Size[c_new] += 1

```

Why does this help?

- `MakeUnionFind(S)`
 - Set `Component[i] = i` for all i
 - Set `Members[i] = [i]`, `Size[i] = 1` for all i

- Find(*i*)
 - Return Component[*i*]
- Union(*i*, *j*)

```

c_old = Component[i]
c_new = Component[j]

for k in Members[c_old]:
    Component[k] = c_new
    Members[c_new].append(k)
    Size[c_new] += 1

```

- Members[*c_old*] allows us to merge Component[*i*] into Component[*j*] in time $O(\text{Size}[c_old])$ rather than $O(n)$
- How can we make use of Size[*c*]
 - Always merge smaller component into the larger one
 - If $\text{Size}[c] < \text{Size}[c']$ re-label *c* as *c'*, else re-label *c'* as *c*
- Individual merge operations can still take time $O(n)$
 - Both $\text{Size}[c]$, $\text{Size}[c']$ could be about $n/2$
 - More careful accounting
- Always merge smaller component into the larger one
- For each *i*, size of Component[*i*] at least doubles each time it is re-labelled
- After *m* Union() operations, at most $2m$ elements have been "touched"
 - Size of Component[*i*] is at most $2m$
- Size of Component[*i*] grows as 1, 2, 4, ..., so *i* changes component at most $\log m$ times
- Over *m* updates
 - At most $2m$ updates are re-labelled
 - Each one at most $O(\log m)$ times
- Overall, *m* Union() operations take time $O(m \cdot \log m)$
- Works out to time $O(\log m)$ per Union() operation
 - Amortized complexity of Union() is $O(\log m)$

Back to Kruskal's Algorithm

- Sort $E = \{e_0, e_1, \dots, e_{m-1}\}$ in ascending order
- MakeUnionFind(*V*) - each vertex *j* is in component *j*

- Adding an edge $e_k = (u, v)$ to the tree
 - Check that `Find(u) != Find(v)`
 - Merge components: `Union(Component[u], Component[v])`
- Tree has $n - 1$ edges, so $O(n)$ `Union()` operations
 - $O(n \log n)$ amortized cost overall
- Sorting E takes $O(m \log m)$
 - Equivalently, $O(m \log m)$, since $m \leq n^2$
- Overall time, $O((m + n) \log n)$

Summary

- Implement Union-Find using arrays/dictionaries `Component`, `Member`, `Size`
 - `MakeUnionFind(S)` is $O(n)$
 - `Find(i)` is $O(1)$
 - Across m operations, amortized complexity of each `Union()` operation is $\log m$
- Can also maintain `Members[k]` as a tree rather than as a list
 - `Union()` becomes $O(1)$

▼ Priority Queues

Dealing with Priorities

Job Scheduler

- A job scheduler maintains a list of pending jobs with their priorities
- When the processor is free, the scheduler picks out the job with maximum priority in the list and schedules it
- New jobs may join the list at any time
- How should the scheduler maintain the list of pending jobs and their priorities?

Priority Queue

- Need to maintain a collection of items with priorities to optimize the following operations
- `delete_max()`
 - Identify and remove item with the highest priority
 - Need not be unique
- `insert()`
 - Add a new item to the collection

Implementing Priority Queues with one dimensional structures

- `delete_max()`
 - Identify and remove item with highest priority
 - Need not be unique
- `insert()`
 - Add a new item to the list

-
- Unsorted list
 - `insert()` is $O(1)$
 - `delete_max()` is $O(n)$
 - Sorted list
 - `delete_max()` is $O(1)$
 - `insert()` is $O(n)$
 - Processing n items requires $O(n^2)$

Moving to 2 dimensions

First Attempt

- Assume N processes enter/leave the queue
- Maintain a $\sqrt{N} \times \sqrt{N}$ array
- Each row is in sorted order

$N = 25$

3	19	23	35	58
12	17	25	43	67
10	13	20		
11	16	28	49	
6	14			

Summary

- 2D $\sqrt{N} \times \sqrt{N}$ array with sorted rows
 - `insert()` is $O(\sqrt{N})$
 - `delete_max()` is $O(\sqrt{N})$
 - Processing N items is $O(N \sqrt{N})$
- Can we do better than this?
- Maintain a special binary tree - **heap**
 - Height $O(\log N)$
 - `insert()` is $O(\log N)$
 - `delete_max()` is $O(\log N)$
 - Processing N items is $O(N \log N)$
- Flexible - need not fix N in advance


▼ Heaps

Priority Queue

- Need to maintain a collection of items with priorities to optimize the following operations
- `delete_max()`
 - Identify and remove item with highest priority
 - Need not be unique
- `insert()`
 - Add a new item to the list
- Maintaining as a list incurs cost $O(N^2)$ across N inserts and deletions
- Using a $\sqrt{N} \times \sqrt{N}$ array reduces the cost to $O(\sqrt{N})$ per operations
 - $O(N\sqrt{N})$ across N inserts and deletions


Binary Trees

- Values are stored as nodes in a rooted tree
- Each node has up to two children
 - Left child and Right child
 - Order is important
- Other than the root, each node has a unique parent
- Leaf node - no children
- Size - number of nodes
- Height - number of levels




Heap

- Binary tree filled level-by-level, left-to-right
- The value at each node is at least as big as the values of its children
 - **max-heap**
- Binary tree on the right is an example of a heap
- Root always has the largest value
 - By induction, because of the **max-heap** property




Non-Examples


No "holes" allowed



Cannot leave a level incomplete



Heap property is violated



Complexity of `insert()`

- Need to walk up from the leaf to the root
 - Height of the tree
- Number of nodes at level 0 is $2^0 = 1$
- If we fill k levels, $2^0 + 2^1 + \dots + 2^{k-1} = 2^k - 1$ nodes
- If we have N nodes, at most $1 + \log \lfloor N \rfloor$ levels
- `insert()` is $O(\log \lfloor N \rfloor)$

`delete_max()`

- Maximum value is always at the root
- After we delete one value, tree shrinks
 - Node to delete is right-most at lowest level
- Move "homeless" value to the root
- Restore the heap property downwards

- Only need to follow a single path down
 - Again $O(\log N)$

Implementation

- Number the nodes top to bottom left right
- Store as a list $H = [h_0, h_1, h_2, \dots, h_9]$
- Children of $H[i]$ are at $H[2 * i + 1]$, $H[2 * i + 2]$
- Parent of $H[i]$ is at $H[(i - 1) // 2]$, for $i > 0$

Building a heap - `heapify()`

- Convert a list $[v_0, v_1, \dots, v_N]$ into a heap
- Simple strategy
 - Start with an empty heap
 - Repeatedly apply `insert(vj)`
 - Total time is $O(N \log N)$

Better `heapify()`

- List $L = [v_0, v_1, \dots, v_N]$
- $mid = \text{len}(L) // 2$, Slice $L[mid:]$ has only leaf nodes
 - Already satisfy the heap condition
- Fix heap property downwards for second last level
- Fix heap property downwards for third last level
- ...
- Fix heap property at level 1
- Fix heap property at the root
- Each time we go up one level, one extra step per node to fix the heap property
- However, number of nodes to fix halves
- Second last level, $\frac{n}{4} \times 1$ steps
- Third last level, $\frac{n}{8} \times 2$ steps
- Fourth last level, $\frac{n}{16} \times 3$ steps
- ...
- Cost turns out to be $O(n)$

Summary

- Heaps are a tree implementation of priority queues

- `insert()` is $O(\log N) O(\log N)$
- `delete_max()` is $O(\log N) O(\log N)$
- `heapify()` builds a heap in $O(N) O(N)$
- Can invert the heap condition
 - Each node is smaller than its children
 - **min-heap**
 - `delete_min()` rather than `delete_max()`

▼ Using Heaps in Algorithms

Priority Queues and Heaps

- Priority Queues support the following operations
 - `insert()`
 - `delete_max()` or `delete_min()`
- Heaps are tree based implementation of priority queues
 - `insert()`, `delete_max()/delete_min()` are both $O(\log n)$
 - `heapify()` builds a heap from a list/array in time $O(n)$
- Heap can be represented as a list/array
 - Simple index arithmetic to find parent and children of a node
- What more do we need to use a heap in an algorithm?

▼ Dijkstra's Algorithm

- Maintain 2 dictionaries with vertices as keys
 - `visited`, initially `False` for all `v`
 - `distance`, initially `infinity` for all `v`
- Set `distance[v]` to 0
- Repeat, untill all the reachable vertices are visited
 - Find unvisited vertex `nextv` with minimum distance
 - Set `visited[nextv]` to `True`
 - Re-compute `distance[v]` for every neighbour `v` of `nextv`

```
def dijkstra(WMat,s):
    (rows,cols,x) = WMat.shape
    infinity = np.max(WMat)*rows+1
    (visited,distance) = ({},{})

    for v in range(rows):
        (visited[v],distance[v]) = (False,infinity)

    distance[s] = 0

    for u in range(rows):
        nextd = min([distance[v] for v in range(rows)
                     if not visited[v]])
        nextvlist = [v for v in range(rows)
                     if (not visited[v]) and distance[v] == nextd]
```

```

if nextvlist == []:
    break

nextv = min(nextvlist)
visited[nextv] = True

for v in range(cols):
    if WMat[nextv,v,0] == 1 and (not visited[v]):
        distance[v] = min(distance[v], distance[nextv] + WMat[nextv,v,1])

return distance

```

Bottleneck

- Find unvisited vertex j with minimum distance
 - Naive implementation requires an $O(n)$ scan
- Maintain unvisited vertices as a min-heap
 - `delete_min()` in $O(\log n)$ time
- But, also need to update distances of the neighbours
 - Unvisited neighbour's distances are inside the min-heap
 - Updating a value is not a basic heap operation

Heap sort

- Start with an un-ordered list
- Build a heap - $O(n)$
- Call `delete_max()` n times to extract elements in descending order - $O(n \cdot \log n)$
- After each `delete_max()`, heap shrinks by 1
- Store maximum value at the end of current heap
- In place $O(n \cdot \log n)$ sort

Summary

- Updating a value in a heap takes $O(\log n)$
- Need to maintain additional pointers to map values to heap positions and vice versa
- With this extended notion of heap, Dijkstra's algorithm complexity improves from $O(n^2)$ to $O((m + n) \cdot \log n)$
- Heaps can also be used to sort a list in place in $O(n \cdot \log n)$


▼ Search Trees

Dynamic Sorted Data

- Sorting is useful for efficient searching
- What if the data is changing dynamically?
 - Items are periodically inserted and deleted
- Insert/delete in a sorted list takes time $O(n)$
- Move to a tree structure, like heaps for priority queues

Binary Search Tree

- For each node with the value v
 - All values in the left sub-tree are $< v$
 - All values in the right sub-tree are $> v$
- No duplicate values



Implementing a Binary Search Tree

- Each node has a value and pointers to its children
- Add a frontier with empty nodes, all fields -
 - Empty tree is single empty node
 - Leaf node points to empty nodes
- Easier to implement operations recursively

▼ The class `Tree`

- Three local fields `value`, `'left'`, `'right'`
- Value `None` for empty value
- Empty tree has all fields `None`
- Left has a non-empty `value` and empty `left` and `right`

```
class Tree:
```

```

# Constructor
def __init__(self, init_val = None):
    self.value = init_val

    if self.value:
        self.left = Tree()
        self.right = Tree()
    else:
        self.left = None
        self.right = None

    return

# Only empty node has value None
def is_empty(self):
    return self.value == None

# Leaf nodes have both children empty
def is_leaf(self):
    return self.value != None and self.left.is_empty() and self.right.is_empty()

```

▼ In-order traversal

- List the left sub-tree, then the current node, then the right sub-tree
- Lists values in sorted order
- Use to print the tree

```

class Tree:
    ...
    # In-order traversal
    def in_order(self):
        if self.is_empty():
            return []
        else:
            return self.left.in_order() + [self.value] + self.right.in_order()

    # Display the tree as a string
    def __str__(self):
        return str(self.in_order())

```

▼ Find a value `v`

- Check value at current node
- If `v` is smaller than the current node, go left
- If `v` is greater than the current node, go right
- Natural generalization of binary search

```

class Tree:
    ...
    # Check if the value v occurs in the tree

```

```
def find(self, v):
    if self.is_empty():
        return False

    if self.value == v:
        return True

    if v < self.value:
        return self.left.find(v)

    if v > self.value:
        return self.right.find(v)
```

▼ Minimum and Maximum

- Minimum is the left most node in the tree
- Maximum is the right most node in the tree

```
class Tree:
    ...
    def min_val(self):
        if self.left.is_empty():
            return self.value
        else:
            return self.left.min_val()

    def max_val(self):
        if self.right.is_empty():
            return self.value
        else:
            return self.right.max_val()
```

▼ Insert a value `v`

- Try to find `v`
- Insert at the position where `find` fails

```
class Tree:
    ...
    def insert(self, v):
        if self.is_empty():
            self.value = v
            self.left = Tree()
            self.right = Tree()

        if self.value == v:
            return

        if v < self.value:
            self.left.insert(v)
```

```
        return

    if v > self.value:
        self.right.insert(v)
        return
```

▼ Delete a value `v`

- If `v` is present, delete
- Leaf node? No problem
- If only one child, promote that sub-tree
- Otherwise, replace `v` with `self.left.max_val()` and delete `self.left.max_val()`
 - `self.left.max_val()` has no right child

```
class Tree:
    ...
    def delete(self, v):
        if self.is_empty():
            return

        if v < self.value:
            self.left.delete(v)
            return

        if v > self.value:
            self.right.delete(v)
            return

        if v == self.value:
            if self.is_leaf():
                self.make_empty()
            elif self.left.is_empty():
                self.copy_right()
            elif self.right.is_empty():
                self.copy_left()
            else:
                self.value = self.left.max_val()
                self.left.delete(self.left.max_val())
            return

    # Convert left node to empty node
    def make_empty(self):
        self.value = None
        self.left = None
        self.right = None
        return

    # Promote left child
    def copy_left(self):
        self.value = self.left.value
        self.right = self.left.right
```

```
self.left = self.left.left
return

# Promote right child
def copy_right(self):
    self.value = self.right.value
    self.left = self.right.left
    self.right = self.right.right
    return
```

Summary

- `find()`, `insert()` and `delete()` all walk down a single path
- Worst-case: height of the tree
- An un-balanced tree with n nodes may have the height $O(n)$
- Balanced trees have height $O(\log n)$
- We will see how to keep a tree balanced to ensure all operations remain $O(\log n)$